

# Cluster 2018

Ciencia de Datos en Ingeniería Industrial

**clase\_02: feature pre-processing**

# AI & Art: Mario Klingemann

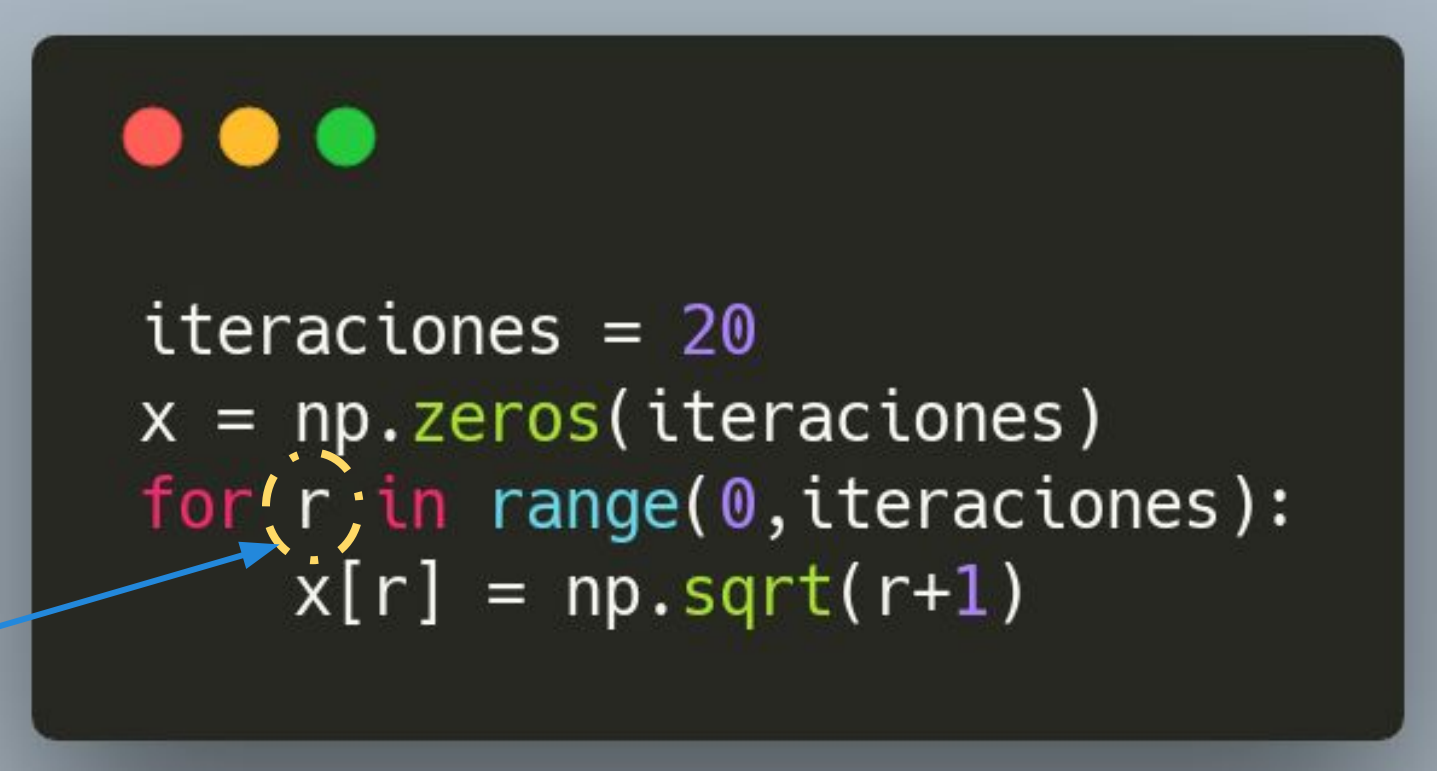


'My Artificial Muse'  
<http://quasimondo.com/>

# agenda\_clase\_02

- Python: For, IF, functions
- Intro Scikit-Learn
- Categorical Variables: Dummies
- Feature Processing
- EDAs
- Practica en clase

# 'for' loops in python



```
iteraciones = 20
x = np.zeros(iteraciones)
for r in range(0, iteraciones):
    x[r] = np.sqrt(r+1)
```

The image shows a code editor window with a dark background and three colored window control buttons (red, yellow, green) at the top left. The code is written in a syntax-highlighted font. The variable 'r' in the 'for' loop is circled with a dashed yellow line, and a blue arrow points from the word 'iterador' to it.

iterador

# 'if' statements in python



```
if x > 1 :  
    x = pd.concat([data1,data2])  
  
else:  
    x = data1
```

# 'if' statements in python



```
if y == "mean":  
    mean = np.mean(data.distance)  
  
elif y == "Preproc":  
    nans = data.isnull().any()  
  
elif y == "std dev":  
    std_dev = np.std(data.distance)
```

# functions in python



```
def dot_product( x1, x2 ):
    "Esta funcion calcula el producto interno de 2 matrices"
    dotprod = np.dot(x1,x2.T)
    return dotprod


a = dot_prod(x_febrero, x_marzo)
```

# Intro scikit-learn





# Intro scikit-learn



## scikit-learn

*Machine Learning in Python*

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

### Classification

Identifying to which category an object belongs to.

**Applications:** Spam detection, Image recognition.

**Algorithms:** SVM, nearest neighbors, random forest, ... — Examples

### Regression

Predicting a continuous-valued attribute associated with an object.

**Applications:** Drug response, Stock prices.

**Algorithms:** SVR, ridge regression, Lasso, ... — Examples

### Clustering

Automatic grouping of similar objects into sets.

**Applications:** Customer segmentation, Grouping experiment outcomes

**Algorithms:** k-Means, spectral clustering, mean-shift, ... — Examples

### Dimensionality reduction

Reducing the number of random variables to consider.

**Applications:** Visualization, Increased efficiency

**Algorithms:** PCA, feature selection, non-negative matrix factorization. — Examples

### Model selection

Comparing, validating and choosing parameters and models.

**Goal:** Improved accuracy via parameter tuning

**Modules:** grid search, cross validation, metrics. — Examples


### Preprocessing

Feature extraction and normalization.

**Application:** Transforming input data such as text for use with machine learning algorithms.

**Modules:** preprocessing, feature extraction. — Examples

# Intro scikit-learn



## scikit-learn

Machine Learning in Python

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

### Classification

Identifying to which category an object belongs to.

**Applications:** Spam detection, Image recognition.

**Algorithms:** SVM, nearest neighbors, random forest, ... — Examples

### Regression

Predicting a continuous-valued attribute associated with an object.

**Applications:** Drug response, Stock prices.

**Algorithms:** SVR, ridge regression, Lasso, ... — Examples

### Clustering

Automatic grouping of similar objects into sets.

**Applications:** Customer segmentation, Grouping experiment outcomes

**Algorithms:** k-Means, spectral clustering, mean-shift, ... — Examples

### Dimensionality reduction

Reducing the number of random variables to consider.

**Applications:** Visualization, Increased efficiency

**Algorithms:** PCA, feature selection, non-negative matrix factorization. — Examples

### Model selection

Comparing, validating and choosing parameters and models.

**Goal:** Improved accuracy via parameter tuning

**Modules:** grid search, cross validation, metrics. — Examples

### Preprocessing

Feature extraction and normalization.

**Application:** Transforming input data such as text for use with machine learning algorithms.

**Modules:** preprocessing, feature extraction. — Examples

# Categorical Variables: Dummies



# Feature Engineering: Categorical Variables (Dummies)

Edad	Altura	Sexo
18	1.70	Masculino
24	1.60	Femenino
30	1.90	Femenino
28	1.5	Masculino



Edad	Altura	Sexo	Masculino	Femenino
18	1.70	Masculino	1	0
24	1.60	Femenino	0	1
30	1.90	Femenino	0	1
28	1.5	Masculino	1	0

# Feature Engineering: Categorical Variables (Dummies)



```
# 1 Creamos un dataframe
raw_data = {'edad': [18, 24, 30, 28],
            'altura': [1.7, 1.6, 1.9, 1.5],
            'sexo': ['masculino', 'femenino', 'femenino', 'masculino']}
data = pd.DataFrame(raw_data, columns = ['edad', 'altura', 'sexo'])

# 2 Creamos un dataframe de variables Dummies para columna "Sexo"
df_sexo = pd.get_dummies(data['sexo'])

# 3 Agregamos estas nuevas variables dummies a nuestro dataframe
df_new = pd.concat([df, df_sex], axis=1)
```



# Feature scaling & normalization



# Feature Engineering: Auto-Scaling & Normalization

En muchas ocasiones las features pueden tener rangos muy distintos. Por ejemplo si utilizamos metros cuadrados y temperatura para caracterizar las condiciones climáticas de un campo la primer variable estará en el rango de decenas de miles y la segunda en decenas, es decir cuando el dominio de las features es distinto. Esto puede generar un problema a la hora de ‘aprender de datos’.

Para eso abordaremos dos estrategias de ingeniería de features:

- Auto-scaling / standarization [1]
- Min-Max normalization [2]

[1] van den Berg, R. A., Hoefsloot, H. C., Westerhuis, J. A., Smilde, A. K., & van der Werf, M. J. (2006). Centering, scaling, and transformations: improving the biological information content of metabolomics data. *BMC genomics*, 7(1), 142.

[2] Jain, Y. K., & Bhandare, S. K. (2011). Min max normalization based data perturbation method for privacy protection. *International Journal of Computer & Communication Technology*, 2(8), 45-50.

# Feature Engineering: Auto-Scaling

El método de “auto-scaling” o “standarization” o “Z-score normalization” asume que cada feature de manera individual responde a una distribución de probabilidad normal y busca estandarizar los valores afectándolos por la media y el desvío standard.

$$x_i' = \frac{(x_i - \mu)}{\sigma}$$

Cada feature después de pre-procesarla quedará con una media = 0 y un desvío standard = 1.



# Feature Engineering: Auto-Scaling



```
from sklearn.preprocessing import StandardScaler
data = [[0, 0], [0, 0], [1, 1], [1, 1]]
scaler = StandardScaler()
print(scaler.fit(data))
    "StandardScaler(copy=True, with_mean=True, with_std=True)"
print(scaler.mean_)
    "[0.5 0.5]"
print(scaler.var_)
    "[0.25 0.25]"
print(scaler.transform(data))
    "[[-1. -1.] [-1. -1.] [ 1.  1.] [ 1.  1.]]"
```


# Feature Engineering: Min-Max normalization

El método de “Min-Max normalization” afecta al valor de la feature en cada sample por el mínimo de la feature y lo divide por el rango entre máximo y mínimo.

$$x'_i = \frac{x_i - x_{min}}{x_{max} - x_{min}}$$

Cada feature después de pre-procesarla quedará un mínimo en 0 y un máximo en 1.

# Feature Engineering: Min-Max normalization



```
from sklearn.preprocessing import MinMaxScaler
data = [[-1, 2], [-0.5, 6], [0, 10], [1, 18]]
scaler = MinMaxScaler()
print(scaler.fit(data))
    "MinMaxScaler(copy=True, feature_range=(0, 1))"
print(scaler.data_max_)
    "[ 1. 18.]"
print(scaler.transform(data))
    "[[0. 0.], [0.25 0.25], [0.5 0.5], [1. 1.]]"
print(scaler.transform([[2, 2]]))
    "[[1.5 0. ]]"
```

# Feature Engineering: Min-Max normalization

**IMPORTANTE:** cuando pre-procesamos las features de un dataset debemos conservar el módulo “scaler” que contiene la información para transformar features. Esto quiere decir que a nuevos datos debemos transformarlos con el scaler ajustado con los datos iniciales y evitar realizar todo el proceso de nuevo con los datos viejos y nuevos.

# A agarrar la PyLA

